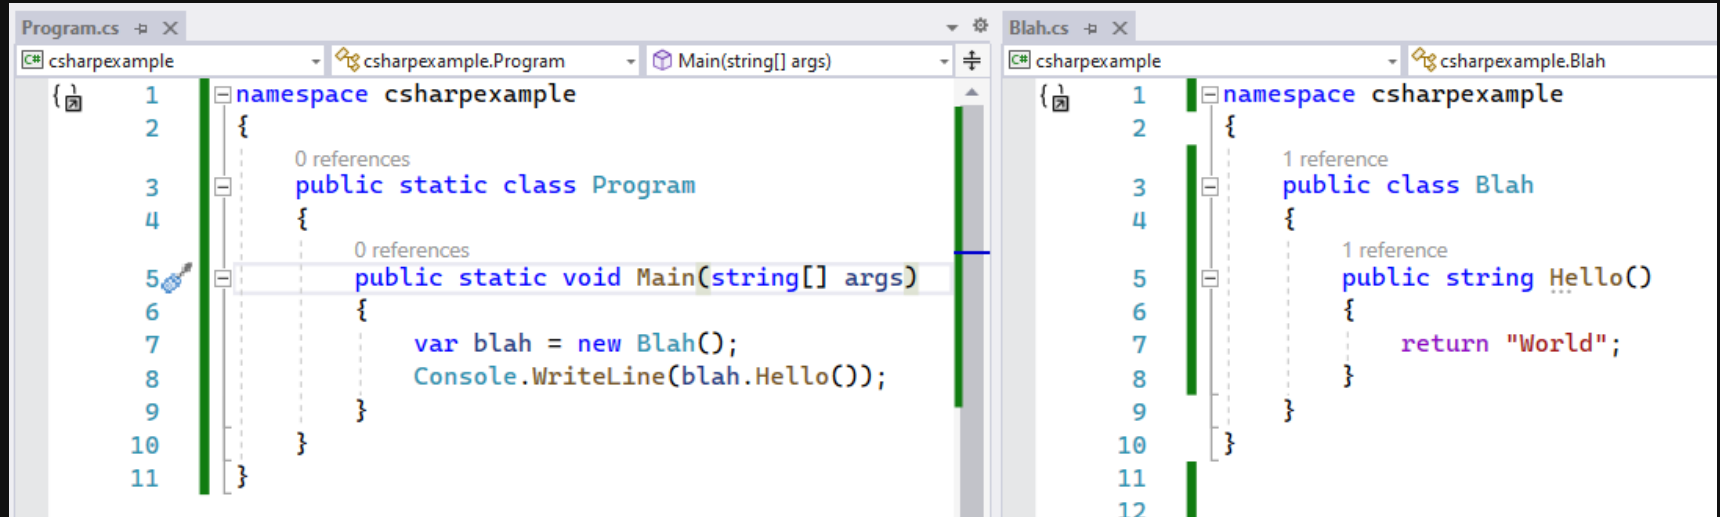


# Translation units & the Preprocessor

# Welcome!

- This is the fundamental “how does C++ compilation work” session
- C++ compiles very differently to most other modern programming languages
- If you don’t understand how things are compiled, you’re going to get confused very quickly
- We’re going to look at how C++ compiles, how it compares with C# and everything that goes on under the hood
- **Aim:** You understand the path your code takes from editor to the compiled binary

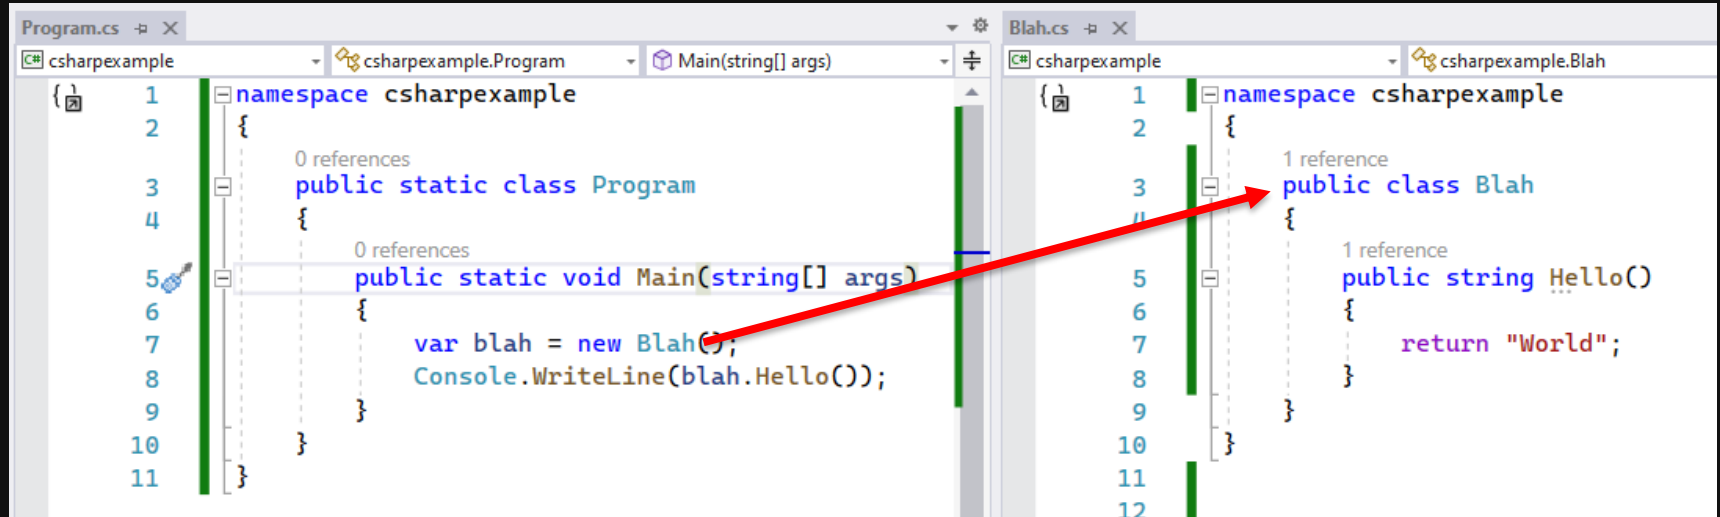
# How does C# compile?



This works because the C# compiler reads all files at once:

```
csc.exe /out:obj\Debug\net6.0\csharpexample.dll Blah.cs Program.cs
```

# How does C# compile?



The screenshot shows two C# files in a project named 'csharpexample'. The left file, Program.cs, contains a namespace 'csharpexample' with a public static class 'Program' and a public static void 'Main(string[] args)' method. Inside 'Main', there is a line 'var blah = new Blah();' and a 'Console.WriteLine(blah.Hello());' call. The right file, Blah.cs, contains the same namespace 'csharpexample' with a public class 'Blah' and a public string 'Hello()' method that returns 'World'. A red arrow points from the 'Blah' reference in Program.cs to the 'Blah' class definition in Blah.cs. The IDE interface includes file tabs at the top, solution explorer on the left, and a breadcrumb trail above the code editor.

```
Program.cs
1 namespace csharpexample
2 {
3     0 references
4     public static class Program
5     {
6         0 references
7         public static void Main(string[] args)
8         {
9             var blah = new Blah();
10            Console.WriteLine(blah.Hello());
11        }
12    }
13 }
```

```
Blah.cs
1 namespace csharpexample
2 {
3     1 reference
4     public class Blah
5     {
6         1 reference
7         public string Hello()
8         {
9             return "World";
10        }
11    }
12 }
```

Because it reads and parses all files at once, it can know what "Blah" refers to in Program.cs!

C++ does not compile this way

# C++ compiles one file at a time

Each file is a “translation unit”

digital

pdp11/70

PAR ERR ADRS ERR RUN PAUSE MASTER USER SUPER KERNEL DATA ADDRESSING 15 18 22

USER D USER I  
SUPER D SUPER I  
KERNEL D KERNEL I  
CONS PHY PROG PHY

DATA PATHS /I ADRS  
BUS REG FPP/CPU  
DISPLAY REGISTER

1972

POWER LOCK  
PARITY HIGH LOW  
OFF

21 20 19 18 17 16 15 14 13 12 11 10 9 8 7 6 5 4 3 2 1 0

LOAD ADRS EXAM DEP CONT ENABLE S INST  
HALT S BUS CYCLE START

# Speculative history time

- Turns out it's quite difficult to find rational for decisions made 50 years ago
- But we can make reasonable guesses as to why the design is the way it is
- So this will partly be a *speculative* history explanation
- Stuff that I could find is from "The Development of the C Language"  
<https://www.bell-labs.com/usr/dmr/www/chist.html>



# Constraints of PDP-7/PDP-11

- C originally written on PDP-7, later ported to PDP-11
- PDP-7 had **8KB memory**. The PDP-11 had **24KB memory**. Total.
- Extremely memory constrained for compilers to operate in.

# What does a compiler need to do?

- Read files
  - Lex (tokenise)
  - Parse
  - Abstract syntax tree
  - Type resolution
  - Emit machine code
- 
- If you had all the source code in one file, you would run into a problem.
  - While the resulting executable would work within memory limitations, there was not enough memory to compile the whole program at once.

# Split up the work

## Compiler

Compile individual .c files to object files

- Read one file
- Lex (tokenise)
- Parse
- Abstract syntax tree
- Type resolution
- Emit machine code as objects

Only need enough memory for compiler to compile each .c file in sequence...


## Linker

Combine object files to get final program

- Read machine code in files
- Stitch machine code together
- Replace placeholder function pointers with actual function locations
- Emit resulting executable

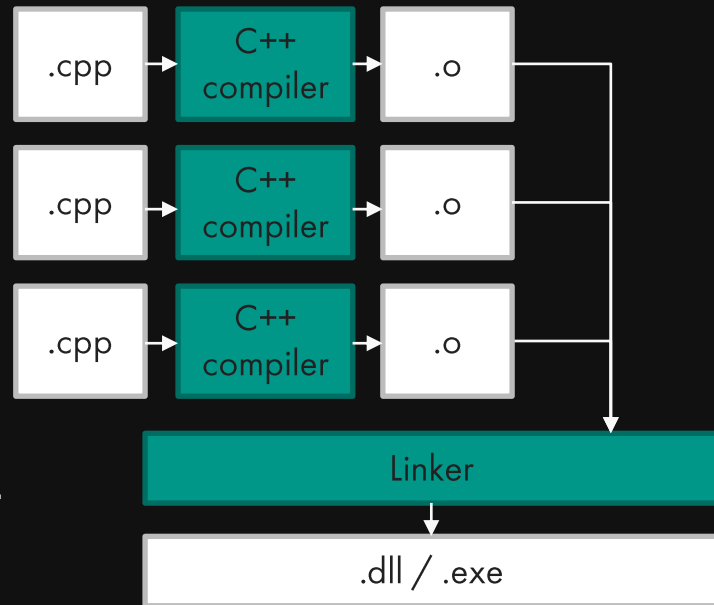
# Comparison again...

## C#



C++ compiler does not know what is in other .cpp files!

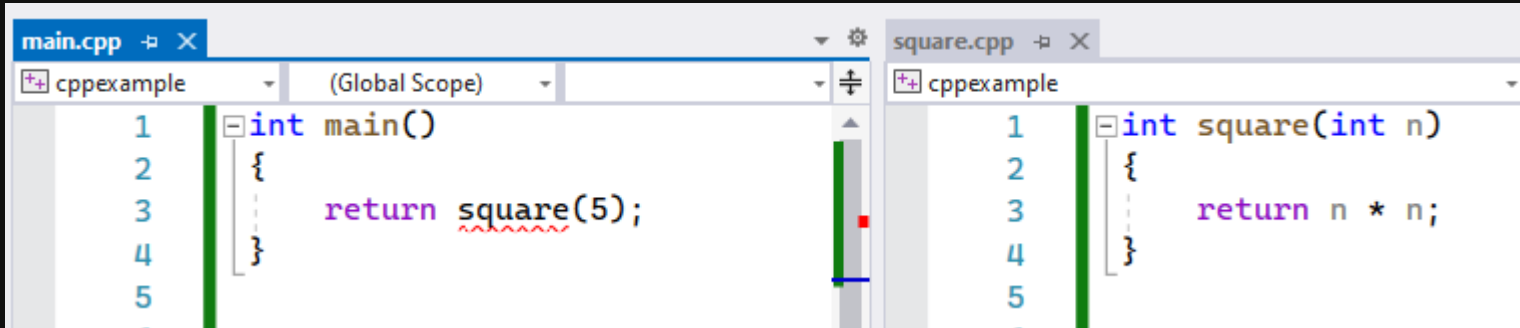
## C++



# How do we reference across files?

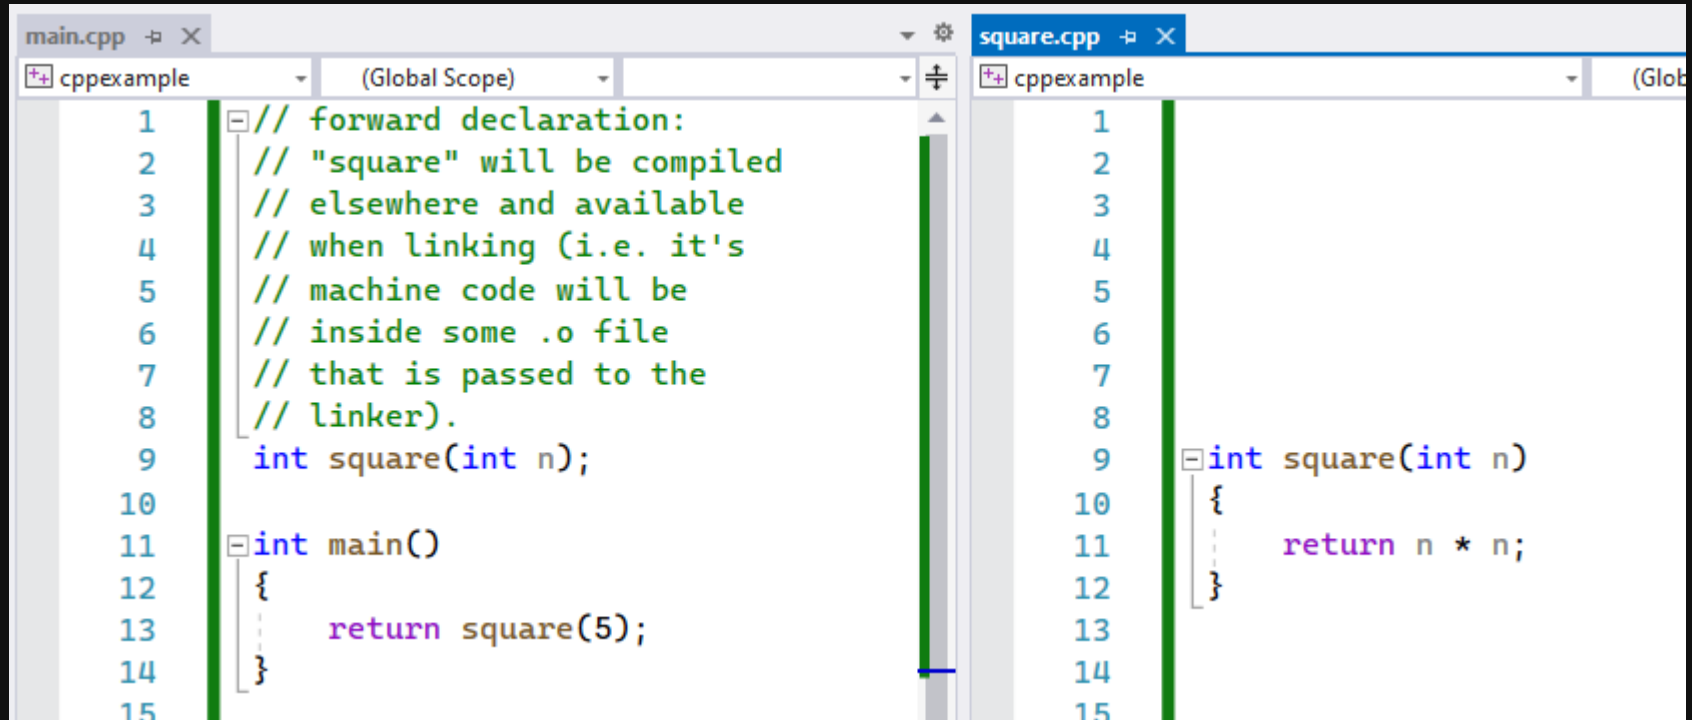
Let's say we want to call `square()` from `main()`.

Given these files will be compiled separately, how can we tell the compiler what "square" looks like in `main.cpp`?



```
main.cpp 1 int main()  
2 {  
3     return square(5);  
4 }  
5  
  
square.cpp 1 int square(int n)  
2 {  
3     return n * n;  
4 }  
5
```

# Forward declarations!



The image shows two side-by-side code editors. The left editor, titled 'main.cpp', contains the following code:

```
1 // forward declaration:
2 // "square" will be compiled
3 // elsewhere and available
4 // when linking (i.e. it's
5 // machine code will be
6 // inside some .o file
7 // that is passed to the
8 // linker).
9 int square(int n);
10
11 int main()
12 {
13     return square(5);
14 }
15
```

The right editor, titled 'square.cpp', contains the following code:

```
1
2
3
4
5
6
7
8
9 int square(int n)
10 {
11     return n * n;
12 }
13
14
15
```

# C++: compiling to executable vs object file

## Input

```
int square(int n)
{
    return n * n;
}

int main()
{
    return square(5);
}
```

Compiler knows address of "square" in this case.

## Straight to .exe file

```
main:
55
push    %rbp
48 89 e5
mov     %rsp,%rbp
bf 05 00 00 00
mov     $0x5,%edi
e8 e3 ff ff ff
call    401106 <_Z6squarei>
90
nop
5d
pop     %rbp
c3
ret
66 2e 0f 1f 84 00 00 00 00 00
cs nopw 0x0(%rax,%rax,1)
```

## Input

```
int square(int n);

int main()
{
    return square(5);
}
```

Compile does not know where "square" is!

Leaves it as 0x00000000.

## To object .o file

```
main:
55
push    %rbp
48 89 e5
mov     %rsp,%rbp
bf 05 00 00 00
mov     $0x5,%edi
e8 00 00 00 00
call    e <main+0xe>
90
nop
5d
pop     %rbp
c3
ret
```

Linker will fill it in later when it has all the object files.

# What does the linker do?

Stitches together object .o files to create the executable.

**Side note:** Also this is where library references are linked in!

## square.o

```
_Z6squarei:
55
push    %rbp
48 89 e5
mov     %rsp,%rbp
89 7d fc
mov     %edi,-0x4(%rbp)
8b 45 fc
mov     -0x4(%rbp),%eax
0f af c0
imul    %eax,%eax
5d
pop     %rbp
c3
ret
```


+

## main.o

```
main:
55
push    %rbp
48 89 e5
mov     %rsp,%rbp
bf 05 00 00 00
mov     $0x5,%edi
e8 00 00 00 00
call    e <main+0xe>
90
nop
5d
pop     %rbp
c3
ret
```

=

main.o will contain extra data (not shown) that tells the linker "please fill in call at XYZ with \_Z6squarei function pointer".



```
_Z6squarei:
55
401106 push    %rbp
48 89 e5
401107 mov     %rsp,%rbp
89 7d fc
40110a mov     %edi,-0x4(%rbp)
8b 45 fc
40110d mov     -0x4(%rbp),%eax
0f af c0
401110 imul    %eax,%eax
5d
401113 pop     %rbp
c3
401114 ret
main:
55
401115 push    %rbp
48 89 e5
401116 mov     %rsp,%rbp
bf 05 00 00 00
401119 mov     $0x5,%edi
e8 e3 ff ff ff
40111e call    401106 <_Z6squarei>
90
401123 nop
5d
401124 pop     %rbp
c3
401125 ret
66 2e 0f 1f 84 00 00 00 00 00
401126 cs nopw 0x0(%rax,%rax,1)
```



# What to do when lots of .cpp files need to forward declare the same function?

- Could forward declare in every file, but tedious.
- Pre-processor created so you can `#include` common code into multiple translation units
- This is a LITERAL TEXT INCLUDE – nothing fancy. You can even run the pre-processor on non-C++ code because it doesn't care about C/C++ code at all.
- So what does the compiler actually see when you compile some normal piece of code?

# What does the pre-processor pass to the compiler?

```
#include <stdio.h>

int main() {
    printf("hello\n");
}
```

## 4 lines of code

**expands to:**

551 lines of code  
that the compiler  
actually parses and  
compiles!

[illegible]

## ***Side note: Seems like a lot of redundant effort...***

- Imagine lots of files that are all including the same headers.
- Each compiler invocation has to parse all of that header code every time.
- Slow and redundant!
- Introducing PCHs: Pre-compiled headers.
- Special compiler mode to compile a set of headers down into a special binary format. Probably a serialized AST (not machine code!)
- Each compiler then loads PCH instead of re-parsing all common headers – speeds up compilation.

# Other things the preprocessor does

- Pre-processor also lets you define macros with `#define FUNC(...) ...`
- Again this is just text processing, so you can emit anything with the pre-processor including things that aren't syntactically valid.
- Macros are very very important to the `UCLASS()` / `UObject` system which we will talk about later.
- But for now important take-away points are...

# THE SUPER IMPORTANT SUMMARY SLIDE

- Each .cpp file is compiled on it's own!
- We use forward declarations so .cpp files can know about things that will be compiled in another translation unit and available to the linker.
- We put forward declarations in headers because we don't want to copy-paste them to every .cpp file manually.
- Compiler actually compiles the fully expanded output from the pre-processor, not just your file!
- All headers and macros are fully expanded before the compiler does any parsing.
- Think about how the #include and macro use is expanding in your file if you're running nearby syntax errors!